

GROK BOT GALAXY

From Customer Research to Service Delivery

Day 2

01	Sales Engineering	Cases, technology, competition
02	Sales	Customer context and follow-up
03	SDRs	Research and prospecting
04	Customer Support	Knowledge, replies, improvement

Livestream: 2026.09.16 23:30 to 2026.09.17 07:53 SGT

Author: Ellen | Created: 17 September 2026

hiellenchuang@gmail.com

READING MAP

Reading map

Start with the full day, then revisit sessions by role. Click a section to jump; click a page number to return here.

03 Overview and timeline

Themes, broadcast order and terminology

06 Sales Engineering

Customer cases, technical answers, competitors and Q&A

09 Live company-building challenge

Game studio pivot, prototypes, collaboration and status

13 Guests and examples

Karen X. Cheng, Matthew Berman, job search and case videos

16 Sales

Research, meetings, CRM, account plans and Routines

19 SDRs

Ideal customers, priorities, research roles and cost

22 Customer Support

Tickets, refunds, knowledge, handoffs and testing

25 Synthesis and application

Shared methods, HR, a working template and sources

How to read these notes

Speaker explanations are the views expressed at the time; live demonstrations are observable operations; editorial synthesis connects sections. Completed demos, ongoing work and unverified results are kept distinct.

Based on the day's video and audio. Official session and broadcast links are on page 28.

01 / OVERVIEW

From research to a customer handoff

Day 2 breaks down sales and support workflows, while the live team changes direction to build a game studio.

Connect the data people actually use

CRM, email, meetings, product usage, knowledge bases and code answer different questions. Once connected, check data freshness and access rights.

Make the next step possible

Research should feed a prospect list, draft, account plan or ticket. Include links, fields, owners and the next action so someone can continue the work.

Teach the workflow your judgment

Specify which customers come first, what makes a good email and when to hand over to a person. Refine those rules using actual results.

Match roles to the work

Some salespeople go straight to specialists; SDRs delegate research through a coordinator; support starts with one Bot. There is no single organization chart.


Design for exceptions

Missing data, expired logins, uncertain answers and policy changes belong in the workflow. Record where work stopped and who handles it next.

Check what actually happened

A draft, written code and a claim that a refund is complete represent different stages. The sessions check tickets, Stripe status and visible results.

Three ideas to connect across the day



This is the reading framework for these notes, not an original presentation slide.

Editorial synthesis of the day's sessions, recordings and guest conversations.

02 / TIMELINE

Day 2 in broadcast order

All times are SGT (UTC+8). The first row starts on 16 September; the rest are on 17 September. Segment boundaries are approximate.

Time	Segment	Focus
23:40 - 00:00	Opening and pivot	Revisit the pop-up event idea; switch to a Bot game viewers can play
00:00 - 00:54	Sales Engineering	Customer case slides, technical answers, competitor research and Q&A
00:54 - 01:43	Game prototype	Start building, landing page and research in parallel; recover from an interruption using shared documents
01:43 - 01:46	Coffee shop case video	Organize family messages and use POS data in operations
01:46 - 02:00	Rules and visuals	Cards, abilities, rounds and the feel of combat
02:09 - 02:34	Karen X. Cheng	Morning newspaper, package alerts and less daily friction
02:41 - 03:30	Live development	Matchmaking, leaderboards, UI, login and team coordination
03:30 - 04:03	Sales	Research, meetings, CRM, account plans and working methods
04:14 - 04:40	Matthew Berman	Content business, family admin, attention and agent roles
04:40 - 04:43	Enterprise case video	Interview clip on Nokia software development and AI use
04:43 - 05:30	Intern and ad production	Job research, cover letters, public work and advertising assets
05:30 - 06:10	SDRs	Prospecting system, priorities, role design and cost
06:10 - 07:00	Game studio progress	Front end, back end, task board, ads, prototypes and sound
07:01 - 07:40	Customer Support	Knowledge, tickets, refunds, human handoffs and evaluation
07:41 - approx. 07:50	Review and wrap-up	Music, game status and next-day work; some features remain unfinished

X records 23:30:16 to 07:53:35 the next day. Rows locate local content; the official full-day end was scheduled for 09:00.

03 / CONCEPTS

A few terms used throughout Day 2

One Bot can have several Routines; the same Skill can also serve different roles.

Term	Meaning in these sessions	Important distinction
SE / Sales Engineer	Helps customers understand technology, demos and product differences	Goes beyond coding: answers must make sense to a customer
SDR / AE	SDRs develop opportunities; AEs move accounts toward a deal	Titles and responsibilities vary; the sessions show the speakers' own practices
ICP	Ideal Customer Profile	Evolves with actual deals, responses and product development
PLG	Product-led Growth	Product usage signals can help decide whom to contact first
Sequence	A staged outreach and follow-up plan	Stop or adjust future actions when a customer's status changes
Routine / Skill	Reusable arrangements for when to work / how to work	Recurring work does not always need another Bot
MCP / Computer use	A tool connection interface / operating a website visually	Use a suitable integration first; visual interaction fills gaps
Trace / Eval	An execution record / checking outputs with test cases	Records support investigation; passing depends on the cases and criteria

Know what each source is for

CRM tracks opportunity status; meeting notes preserve what a customer said; the knowledge base defines product and policy information. Bots can search across systems, but conflicts between sources still need to be resolved.

A reading aid based on classroom usage, not a complete product specification.

04 / SALES ENGINEERING

Three roles for three customer questions

Amrita Venkatraman | Approx. 00:00 - 00:54

The session separates presales work into customer examples, actual product behavior and competitive differences. Mimi uses case-study slides from everyday work; Sherlock and Serena operate in the FlyLo demo environment.

Bot	Task and sources	Deliverable
Mimi	Read public customer articles and existing cases; reuse the master deck template	Case-study slides for customer presentations
Sherlock	Read FlyLo repositories, including front end and back end, to understand actual capabilities	Technically grounded answers suitable for customers
Serena Williams	Explore competitor websites, workflows and new capabilities	Competitive comparisons, differences and suggested follow-up research

Mimi: specify the structure, then add content

- Slides follow a consistent Problem, Solution, Impact and quote structure, retaining company logos and a shared visual layout.
- The live task uses a public Salesforce article, then assigns a Grab article. Jellyfish is an existing case shown during the demo.
- The Bot creates slides in its own browser. When login fails, the speaker takes over to resolve it before letting the Bot continue.
- The speaker reviews the generated Salesforce slides. The Grab task is still running during the segment; both should not be described as fully reviewed.

What could become a Routine

Regularly find new public customer stories, assess whether they belong in the master deck, then use a fixed template. Results reported in an article still need their original source; reformatting does not independently verify them.

Based on the Sales Engineering recording, audio and official session page.

05 / SALES ENGINEERING

Connect technical answers with competitor research

FlyLo is a classroom demo product. The capabilities and discussion below belong to that example.


Sherlock: inspect the implementation, then explain it

The speaker asks what happens when two people book the last available place at the same time. After reading the code, Sherlock explains the database safeguards and translates them into a customer-facing answer: checkout holds the place for a period, and a later buyer encounters a conflict if someone else has already secured it.

The answer mentions a hold of about 10 minutes. This describes the FlyLo example, not a real airline policy. A Bot's explanation is also not evidence that a concurrency test passed.

Serena: understand both the competitor and your product

- After receiving a competitor-research task, Serena asks Sherlock which search, fare, seat and baggage flows FlyLo currently supports.
- The speaker names Southwest and Spirit, then narrows the question from general booking flows to capabilities such as AI travel assistants.
- Competitor claims, observed website behavior and FlyLo code are three different kinds of evidence to compare.



Use a group to bring information together

The speaker brings both Bots into one discussion to ask which differences matter and which would be easier to implement. The deliverable at this stage is analysis and recommendations, not a released product feature.

Based on operations around 00:15 - 00:28 and subsequent review.

06 / SALES ENGINEERING

Turn a demonstration into a reusable method

Skills, additional roles and questions from the audience.

- The speaker takes over the Bot's computer, demonstrates how to read technical updates or engineering blogs and identify AI changes, then asks it to apply the method to other competitors.
- The task grows from reading one article to knowing how to find similar information. Later results still need checks for sources, freshness and actual use of the taught method.
- The speaker asks for new teammates beyond the existing roles. Battle Card Blair, Demo Drake and AI Radar appear with responsibilities for comparison materials, demos and technology research.

Question	Answer or limitation shown
What if there is no MCP?	Use the Bot's browser to interact with websites. This does not guarantee successful login or passage through bot detection.
API or visual interaction?	An efficient working integration was usually preferred. The speaker's view of future visual-interaction speed is a prediction.
How are important actions restricted?	Enterprise settings can constrain websites and actions. Production deployment and internal-system use can require permission first.
Windows- or Mac-only software?	The speaker says the Bot VM was Linux. Software for a specific operating system cannot be assumed to install directly.
Why are tokens consumed?	Conversation, tool use and related work may incur usage. Reducing unnecessary reporting is one possible improvement.

A method to reuse

Provide an example you approve, explain its format, sources and limits, then check whether the Bot can repeat the same type of work. Team consistency comes from shared methods and traceable evidence, not similar-sounding role names.

Based on the screen demonstrations, operations and Q&A around 00:29 - 00:54.


07 / LIVE BUILD

From a pop-up event to a game studio

The three-day company-building challenge with Matt Palmer, Lauren Tan and Roshan Sadanani.

Day 1 focused on an experiential pop-up event. Guest feedback on venues, production and practical execution prompted the team to reconsider what could fit into three days. On Day 2, they pivot to a game studio, aiming for something both online viewers and people at the venue can play.

The smallest game: turn Bot templates into characters



- Templates can share Skills, Routines and selected general memories. The team wants to turn these settings into characters, rather than let Bots arbitrarily decide each match's rules.
- Early discussion includes Charisma, Dexterity and Intelligence, a fixed pool of points and character traits. Names and numbers continue to change during implementation.
- The same template should produce an understandable, consistent character foundation, before adding appearance, rarity and variation.
- First establish team selection, combat and result review; then consider more complex arenas, upgrades, leaderboards and sharing.

The business direction changes too

The studio is the business; the game is its first product. The team also discusses distribution, support, community and revenue. Paid cosmetics, venues or arenas, and ads are candidate models. No revenue is established during the day.

Based on opening audio and later recordings. Ideas are described as ideas, not delivered products.

08 / LIVE BUILD

Find the fun before designing the full architecture

A prototype gives the team something concrete to discuss about play and presentation.

Practice shown	What it helps with	Still to check
Read the shared repo plan after an interruption	Rebuild context when a new Agent lacks the full earlier conversation	Understands the current scope instead of outdated decisions
Restate rules; get human confirmation before editing	Expose misunderstandings before work goes in the wrong direction	Includes constraints and acceptance criteria
Potato Mode / lightweight HTML prototype	Compare mechanics and interfaces at low cost	Whether it is fun, and the cost of temporary simplifications
Multiple models and design variants	See different approaches, then select and combine good ideas	Whether all candidates have finished before comparison
Sliders and debug views	Change stats and rules immediately and observe the effect	Numbers match the display and rules remain consistent
Play a round, then give feedback	Find real issues with flow, outcomes and animation	Unexpected changes in reset behavior, ordering or results

The prototype deliberately skips login, uses in-memory data and seeded Bots, and offers three switchable UI variants in HTML, CSS and JavaScript. Rules, values and interactions come before the production architecture. A debug panel can simulate many games to explore the effect of ability bonuses.

Around 01:57, the team decides to try single-round combat first. Early collections, replay IDs, Elo and leaderboards include mock data. An interactive prototype does not establish completed multiplayer matchmaking, persistent storage or game balance.

Say exactly what the feedback concerns

An Agent misunderstands design requests and makes the visuals too elaborate. The speakers suggest turning recurring mistakes into a Skill so future work can reuse the correction, beyond a one-time manual fix.

Based on live development around 00:54 - 02:00 and 02:41 - 03:30.

09 / LIVE BUILD

A shared foundation for parallel development

Build features while closing gaps between design, code and coordination.

Give everyone the same product to work on

- The team wants an agreed prototype on main before people develop features independently, with an engineering Bot helping track dependencies and PRs.
- A whiteboard breaks the UI into characters, team selection, matchmaking, combat and results. When there is a lot of data, make the next action clear instead of displaying every number at once.
- The prototype later moves into Next.js / React, separating front end, server state and API contracts. Interfaces can be explored in parallel; shared state needs common rules.
- X login works in the wrap-up, but Bot import fails and the leaderboard is wrong. Login, matchmaking and complete combat require separate checks; early mock lists are not real players.

Connect task tracking to actual changes

Notion tracks tasks, Slack shares decisions, and the engineering Bot follows the repo and PRs. Some cards move to Done while rebasing and merging remain. Submitted, merged and verified describe different stages.

Shared item	Question it should answer
Task board	Who is working, what is blocked and when can it be reviewed?
PR / branch	What changed, can it coexist with other work and who checks it?
Shared document	Which rules apply now, and which decisions have changed?
Working version	Can a user actually complete the whole flow?

Separate implementation from verification

The speakers introduce long-running coordination and shared context in Cursor Projects, then show Full Autopilot in Potato Mode: implementation agents build features; verification agents run the code and operate the UI. This describes the demonstrated workflow, not proof that all PRs passed or merged.

Based on recordings around 02:41 - 03:30, 06:10 - 07:00 and the wrap-up after 07:41.

10 / LIVE BUILD

Visuals, distribution and what remains unfinished

Moving from a visible interface toward a game people want to play.

Workstream	What appears during the day	How to read the status
Visuals and animation	Show Imagine + FAL background removal; assign matching ability icons	Assets and prototypes appear; new icons still need delivery checks
Landing page	Introduce the product, avoid a login-only first impression and add sharing details	Shown and revised; every entry point was not verified
Ad assets	Reuse React components in Remotion; adjust 1:1, 9:16 and 16:9 formats	Assets appear; cropping and content still need work
Advertising business	Discuss lobby sponsorship, advertisers, campaigns and creative data	Business and technical exploration; no evidence of ad sales
Music and sound	Start an audio Bot around 06:58; try Strudel and Suno	A sample plays around 07:42; full game integration is not confirmed
The game	Hosts say the back end is connected and demonstrate X login; Bot import fails	Login is shown; import, leaderboard and complete combat still need fixes and checks

Ask Bots to review their own workflow

Dr. Eggbot creates game-design and PlayTester roles so testers can operate the game. The team also asks it to inspect work conversations for delays from manual merges and waiting. Lauren still wants to participate in prototype trade-offs. Working features, enjoyable play and more autonomous progress are separate questions.

An accurate description of the day's result

The team moves from a pivot to a game prototype and engine, with several active workstreams. Deployment, naming and audience play are left for later; the advertising MVP remains untested. The process shows rapid experimentation and integration work, not a fully successful game business.

Based on live development and the wrap-up around 07:41 - 07:50; demonstrations, plans and unchecked work remain distinct.

11 / GUESTS

Karen X. Cheng: stay informed with less phone time

Approx. 02:09 - 02:34 | Also includes the coffee shop video around 01:43.

Morning Newspaper: turn scattered information into paper

Karen starts with a simple need: know whether anything urgent happened overnight without picking up a phone and getting drawn into scrolling. A Bot gathers calendar items, email, packages, weather and selected content into a personal newspaper, scheduled to print overnight or early in the morning. She shows the printed result.

- She iterates on PDF / CSS layout, type and assets, then turns them into a template. Others can configure their name, location, photo and interests. Reusable assets reduce repeated layout work.
- She also describes package tracking: check more often as delivery approaches, then send an alert to the split-flap display at home.
- Cursor requires more involvement in implementation details. She feels the Bot handles more setup and troubleshooting, making previously burdensome ideas worth trying.
- Expired sessions and repeated logins remain points of friction. Connecting a physical printer also shows that a workflow can cause real-world actions.

Start with a small task you would rather not keep doing

Ask what you must watch for every day but often miss. Identify the inputs, how to organize them, where to show the result and when to start. That gives you a workflow to test.

Coffee shop case: bring organized information into operations

The coffee shop operator describes two uses: organizing school emails and schedules for the family, and connecting POS data to inspect product sales. These are the interviewee's reported experiences. The clip provides neither complete financial data nor an independent impact evaluation.

Based on the guest conversation and case video; the Morning Newspaper creator was checked against the official Marketplace.

12 / GUESTS

Matthew Berman: save attention for creative work

Approx. 04:14 - 04:40 | Content business, family administration and managing several Bots.

Context	Work given to Bots	The person's role
Content business	Sponsor correspondence, meeting notes, proposals, follow-up and daily administration	Keep decisions on creativity, topics and content direction
Family plans	Read school emails; organize events and schedule conflicts	Confirm preferences and important trade-offs
Secondhand items	Organize unused items, research prices, prepare listings and replies	Handle login, important terms and physical handover
Bills and subscriptions	Read plan details, compare fees and identify possible savings	Check calculations and suitability
Work context	Find the last conversation, responsible person and past access blockers	Use the recovered context to decide and proceed

The burden goes beyond switching tools

Reconstructing context also takes time: what someone said last time, where a document is, and who owns a decision. Connecting this information can make small tasks worth doing even when they previously seemed to offer too little return.

Coordination can help, but can also create more management

The conversation discusses a Chief of Staff to organize work and Dr. Eggbot to create roles. But if a person still visits every Bot to inspect details, they become a manager with too many reports. Summaries, exception reporting and clear responsibility matter more than the number of roles.

The next video: Nokia's engineering example

An enterprise interviewee describes AI use for large-codebase analysis, decomposing monolithic systems and root-cause investigation. This is a reported experience, not a complete engineering validation repeated live. These notes omit precise efficiency figures that were not independently checked.

Based on the guest interview; personal experience is not a general guarantee.

13 / GUESTS

Job applications need research and feedback too

Shardul Marathe | Approx. 04:43 - 05:28, interleaved with ad production.

The workflow shown

- Recruiter Finder researches recruiting contacts. A job-search role gathers descriptions from public vacancies and hiring systems, passes them to Cover Letter Writer, and Critic reviews the result.
- Cover letters use a resume and a file of personal experiences. The guest specifically asks the Bot not to invent experience. The aim is to match the role, beyond replacing company names in bulk.
- Critic points out vague content, tone problems and sentences that do not address the role. Repeated feedback helps it learn the guest's writing preferences.
- A contact CSV, cover-letter drafts and plans to connect email represent different completion states. Private alumni lists and contact details are not reproduced here.

The hosts' broader career advice

An application should help a team see what work you can do. Alongside applying, build projects, share work publicly and join relevant communities so people with similar interests understand your abilities. The conversation does not establish that using a Bot guarantees a job.

Generic draft	A more specific direction
I am passionate about AI	Explain the problem the organization is working on and where you could contribute
I have many relevant skills	Link to a real project, your responsibility and the result
Make this letter better	Specify the reader, purpose, tone, unwanted phrasing and successful examples

Keep the lesson beyond this one letter

When the same mistake recurs, save the reason, a good example and a counterexample. This also applies to proposals, research summaries, sales emails and internal documents.

Based on the interview and demos; the name was checked on the guest's profile shown around 05:27.

14 / SALES

Keep customer context close to the work

Krista Letz and Mark Wright | Approx. 03:30 - 04:03

The session focuses on reducing search, organization and repetitive updates. Bots need access to the data people actually use to connect research, meeting preparation, customer signals and follow-up actions.

Role	Sources and tasks	Deliverable
Olive / Chief of Staff	Calendar, Slack, email and daily priorities	Meeting preparation, inbox summaries and editable email drafts
PG / Outbound Prospecting	CRM accounts, public information, interviews and personal updates	Prioritized prospects, grounded outreach angles and drafts
Echo	Meeting records from Granola, Gong and similar tools	Updated slides covering customer use cases, outcomes and next steps
Customer Expert	Notion account plans, Slack, usage and product changes	Renewal preparation, stakeholders, signals and follow-up items
Engineer	Repositories and referenceable examples	Customer-ready technical explanations and instructions

Personalization takes more than changing a company name

Krista looks for what the recipient actually cares about: issues an executive discusses, product direction, interviews and recent posts. She connects that context to useful value. She asks Bots to go beyond returning a video link: read its content, extract useful information and prepare the next step.

Give the tone a source too


Provide Gmail, Slack or other writing examples you approve and explain what does not sound like you. Whether to send still depends on the actual draft, its evidence and its reader.

Based on the Sales recording, official session page and What we learned slide.

15 / SALES

After the meeting, move information into action

From hearing customer needs to updating slides, CRM and account plans.



Echo: this is not a promise of continuous live streaming

The Q&A clarifies that the speaker stops Granola recording so Echo can access the resulting transcript and update the slides. A short pause during a meeting can allow this. Gong data may require waiting until afterward. The demonstration should not be described as instant synchronization without delay.

CRM: first define what must be recorded

- The speaker specifies initials, date, the customer's intended outcomes and next steps, asking the Bot to fill these from email, Slack and meeting notes.
- A person reviews the proposed update before it moves into Salesforce. A completed draft and a completed CRM write need separate confirmation.
- Account plans in Notion maintain projects, stakeholders, usage signals and renewal context.

Product updates can reconnect with earlier promises

When a feature appears in the changelog, Customer Expert can identify customers who previously requested it and remind sales to follow up. This ties the update to an earlier customer problem instead of broadcasting only that a feature has launched.

Reusable handoff fields

What the customer said / source and date / what we promised / who does the next step / when to reply / what still needs confirmation. This is an editorial template, not a verbatim classroom instruction.

Based on demonstrations around 03:38 - 03:47 and the Echo Q&A around 04:01.

16 / SALES

Three takeaways and different ways to divide work

The Sales session wrap-up and Q&A.

Connect everyday tools

Bring CRM, email, calendar, Slack, Notion and meeting tools into the workflow. Sources retain distinct purposes; everything need not be moved into one place first.

Start with one clear job per role

Meeting prep, prospect research, inbox work or technical customer support: give a clear task, examples and feedback before expanding.

Use Routines for recurring work

Schedule morning preparation, overnight drafts and regular research. Adjust frequency to freshness, noise and usage.

Share templates, then adapt

Start from an experienced colleague's template. Keep a common foundation, then add your own customers, voice and working method.

Q&A topic	Key point
Must work go through a Chief of Staff?	No. Krista often contacts specialist Bots directly; Mark uses a coordinator and Routines more often. Both adapt to their work.
One Bot per customer?	A few strategic accounts may merit deeper retained context. Large account sets need a different approach; assess maintenance cost and data boundaries.
How can usage be controlled?	Reduce unnecessary Routine frequency, prefer a mature MCP when suitable, and avoid sending every small task through several layers of conversation.
How do teams collaborate?	The demo still uses existing channels such as Slack. Bots can track key information, but this does not establish a new cross-organization system.

Based on the What we learned slide and Q&A around 03:50 - 04:03.

17 / SDRS

Set priorities before scaling prospecting

Simon Lackowski | Approx. 05:30 - 06:10

Simon describes SDR work as an ongoing system: find suitable companies and people, enrich the data, understand the customer, write messages, plan follow-up and adjust to opportunity status. The aim is to know what to do first each day.

Work	Inputs	Judgment required
Prospecting	ICP, deals, opportunity data and public research	Which titles and companies are likely to respond or advance?
Account Research	Usage, funding, vacancies, technology and organization information	Why contact them now, and what problem might they be addressing?
Sequencing	Continuously updated lists, AE meetings and CRM status	What happens next, when, and when should it stop?
Drafting Copy	Your emails, positive responses and recent successful examples	How can your own voice address what matters to this person?

The ICP should keep evolving

He stores the ideal customer profile as an updateable Skill, refining it with opportunities that advance, customer priorities and new product capabilities. A previously lost prospect may become worth revisiting when a missing feature is added.

Do not simply imitate every old email

Start with your own accounts and outgoing messages, then identify which received positive responses. Give recent effective examples more weight. Continue reviewing individual drafts so the Bot does more than swap company and person names.

Based on the full transcript of the local SDR recording, slides and official session page.

18 / SDRS

Simon's roles and daily workflow

This is the speaker's own setup, not a required organization for every SDR.

Role / group	Responsibility
Simon Bot / Chief of Staff	Main conversation entry point; knows the other roles and their handoffs
Shakespeare	Learn the speaker's writing style and prepare email messages
Web Search + Simon Soldiers	Distribute company research and collect the results
PLG / Usage / Voice of Customer	Combine CRM, usage, customer conversations and lost-deal reasons
Enrichment and Company Research	Use tools such as Amplemarket and Sumble for contact, technology and organization clues

Before the workday starts

- The speaker sets a daily list of 50 prospects, picks the 5 needing immediate attention, and schedules the rest. This is his setup, not evidence of 50 qualified opportunities produced every day.
- Calendar and inbox information help decide which messages need a reply now and which tasks can wait until the afternoon.
- Scan new signups, usage or content downloads on workdays and act while the context is fresh.
- Changes in Salesforce stage can stop or adjust a sequence, avoiding inappropriate prospecting messages after an opportunity has progressed.

A CSV is a tracking artifact, not a completion metric

The screen shows lists, queues, Gmail draft IDs and contact information. These make drafts and recipients traceable. Delivery, replies and opportunity quality require separate checks.

Based on the SDR screen demo and audio; exact counts illustrate the speaker's configuration only.

19 / SDRS

Parallel research, shared reporting, purposeful roles

The SDR wrap-up and Q&A.

Why have a separate research group?

Web Search assigns accounts to research Bots that need little prior context and collects their results in a group. Simon says this makes each subtask easier to inspect and research volume easier to adjust, without routing every small task directly through the top-level coordinator.

Three lessons from the session

- **Connect the whole workflow.** Move from prospecting and enrichment to context, drafts and follow-up. Teaching and revision take time upfront before they can reduce repeated prompting.
- **Design roles deliberately.** Create roles through the main coordinator so it knows their purpose. Adding Bots should not leave the person managing more fragmented conversations.
- **Ask whether a new role is necessary.** One email Bot may also handle the inbox; recurring work may need only a Routine. Different platforms or parallel tasks offer stronger reasons to split roles.

Cost and getting started	The speaker's answer
What does this cost each day?	It depends on meetings, data volume, new-list needs, existing context and scale. No universal daily price was given.
Repeat research every day?	Compare small daily batches with larger weekly batches, considering freshness, usage and actual work patterns.
Unsure how to split the work?	Describe your current work, pain points and desired result. Ask the Bot to propose a setup, then refine it.

Editorial note: self-ratings are only a reference

A Bot can score its own email, but liking its writing is not an accuracy or response rate. Still sample-check sources, recipients, tone and real outcomes.

Based on the latter part of the complete SDR audio, organized into readable working principles.

20 / CUSTOMER SUPPORT

Read first, draft next, then consider direct replies

David Gan | Approx. 07:01 - 07:40

David recommends introducing automation in stages: first read tickets and identify the underlying issue, then leave drafts in tickets for people to review. Only later, based on test results and permission design, allow direct replies for selected situations.

Use	Work involved
Reply to tickets	Find answers in public documentation and company policy; decide whether to respond or hand over
Investigate first	Organize the customer issue, information already checked and potentially relevant context
Flag important cases	Identify locked accounts, SLA conditions and churn risks; notify the appropriate team
Answer internal questions	Use the same knowledge base for colleagues while separating public and internal information by audience
Improve the system	Review tickets and execution records to find missing documentation and process changes

Four roles in the demo environment

Build + Reply

Build establishes connections and initial settings. Reply uses tickets, knowledge and procedures to prepare answers or perform permitted actions.

Alert + Tune

Alert sends important notifications. Tune identifies gaps, proposes changes and updates knowledge after approval.

You do not need to copy all four roles at the start

The speaker explicitly recommends learning one workflow with one Bot first. Split responsibilities when the scope grows or parallel work becomes useful.

Based on the complete support recording transcript, slides and official session page.

21 / CUSTOMER SUPPORT

Four test cases, four response paths

The FlyLo demo uses Notion, Plain, Stripe and Slack in a classroom test environment.

Case	Behavior shown	What can be confirmed
Forgotten password	Read public auth documentation; log the issue, source and confidence; provide reset steps	A response with source comparison is demonstrated
No Okta / SSO answer in the knowledge base	Recognize insufficient information and hand off; notify Alert and flag an important locked-out customer in Slack	Stopping automated answers and sending a notification are demonstrated
Two refund requests	Apply subscription timing and a 14-day demo policy; approve one and reject the other	The speaker refreshes Stripe: one subscription is canceled and refunded; the other account is unchanged
Can a Wi-Fi pass be shared?	Identify the missing FAQ answer; after approval, add a no-simultaneous-sharing rule and run the reply again	Knowledge is added first, then the answer cites the new section

Use the same knowledge differently for different readers

External replies do not paste the complete internal SOP. Authorized colleagues can receive the refund procedure in Slack. The live Slack response waits for and requests sending permission before the reply appears.

Knowledge-base changes also need review

An incorrect rule saved into the knowledge base can affect every similar question afterward. The improvement workflow asks the speaker first. The ability to improve does not mean permission to invent company policy.

Based on the recording around 07:13 - 07:29 and the full transcript. Refund terms apply only to the demo policy.

22 / CUSTOMER SUPPORT

Trace problems and test changes before using them

Support Q&A: records, costs, permissions and knowledge updates.

A trace should explain what actually happened

The speaker describes recording execution time, documents consulted, the information ultimately used and the answer, even for a dry run or an internal ticket note. He mentions Supabase Postgres tables for traces, eval definitions and eval runs, reviewed by an improvement role. The individual database writes are not verified in this segment.

Knowledge changes can use branches and PRs

Notion makes the demo easy to follow. For stricter change control, the speaker suggests storing knowledge in a version-controlled environment and using branches, PRs, code owners and approval before merging. Run test cases on the branch before applying the change to check the expected behavior.

Question	Key point
How should costs be estimated?	The speaker estimates US\$1-2 for his medium-to-complex tickets, and about US\$0.20 for simple types after grouping and scripting. These are personal experiments, not official prices or guaranteed rates.
How can wasted searching be reduced?	Provide a specific ticket ID, necessary data and a defined set to process. This usually takes less searching than simply saying "reply to Alex."
How can nontechnical colleagues use it?	Set boundaries in templates, connection permissions and review procedures. Natural-language reminders alone do not protect production data.
Which support work should come first?	Start with frequent, repetitive types backed by clear knowledge. Improve documents and SOPs, test, then expand gradually.
What about phone support?	The speaker mentions phone experiments, but a complete phone-support workflow is neither shown nor verified here.

Based on Q&A around 07:29 - 07:40. The speaker describes database records; this segment does not fully demonstrate them.

23 / SYNTHESIS

The same questions across different kinds of work

The following is an editorial synthesis across sessions.

Shared question	Sales / SDRs	Support / product development
Where does context come from?	CRM, meetings, usage, public information and earlier correspondence	Tickets, policies, code, shared plans and user feedback
What determines the next step?	ICP fit, fresh signals, opportunity stage and importance	Sufficient knowledge, impact, dependencies and test results
How can good work be repeated?	Message examples, research Skills, fixed fields and Routines	SOPs, knowledge versions, regression cases and repeatable prototypes
When is a person needed?	Commitments, relationships, external messages and resource trade-offs	Policy exceptions, important writes, knowledge changes and product direction
What proves completion?	CRM changes, reviewable drafts, actual replies and opportunity progress	Ticket outcomes, transaction status, tests and a complete usable flow

Let roles evolve with the work

Krista goes straight to specialists, Simon delegates research through a coordinator, and David recommends starting with one workflow. These approaches fit different information volumes, parallel tasks, shared context and management costs.

A useful question before adding a Bot

If this Bot were removed, what necessary result would disappear? If it merely passes messages along, simplify the setup. Independent research, specialist judgment or parallel execution offers a clearer reason to keep a separate role.

Draws together Sales Engineering, Sales, SDRs, Customer Support and live development.

24 / APPLICATION

Where these methods could fit in HR

Editorial extension. These are not HR outcomes demonstrated by the Day 2 speakers.

HR context	An initial deliverable	What a person should confirm
Employee policy questions	Find applicable clauses, versions and sources; prepare a reply draft	Exceptions, sensitive cases and final decisions
Recruiting research	Organize public vacancies, labor-market context and role differences with dates and sources	Qualification judgments, fairness, candidate contact and personal-data use
Interview and meeting follow-up	Extract next steps, deadlines and owners from authorized records	Accuracy, wording of subjective assessments and external commitments
Onboarding support	Bring together schedules, document locations, tasks and blockers	System access, exceptions and individual needs
Reports and data checks	Check fields, omissions, sources and discrepancies; flag items for investigation	Data use, pay and employment decisions, and production-system writes

Start with work whose results are easy to inspect

A useful starting point is classifying employee questions and attaching relevant policy, while leaving employment decisions to people. Compare with existing human answers, then check sources, versions, missed cases and situations that require handoff.

Run a small scope first

Choose a small set of traceable cases. Let the Bot organize and draft, with human review before use. Record error types and corrections. Once the workflow is stable, consider increasing frequency, sources or write permissions.

Workflow-design examples derived from the methods discussed in the sessions.

25 / WORKING TEMPLATE

Define one workflow clearly

Use this template to organize your first Bot task.

Item	What to specify
Goal	Which burden should be reduced? What is the final deliverable?
Trigger	A person's request, a schedule or a changed event?
Data	Which systems, fields, documents and versions can be used?
Steps	What is checked first, how are decisions made and when do other roles take over?
Output	A draft, table, task, PR or actual system change?
Boundaries	What is read-only, what can be written and what needs permission first?
Exceptions	What happens with missing data, low confidence, expired login or conflicting sources?
Acceptance	Which test cases and real results establish completion?
Feedback	Who corrects it, where is the correction stored and how is it reused?
Frequency and cost	How often is fresh data needed? Must everything be checked again each time?

Four questions after completing a task

- Can the result be opened, inspected and traced to its original sources?
- Does the next person know what to pick up and what is still missing?
- Do important exceptions actually stop or hand over, rather than continue guessing?
- Has the correction been saved in the appropriate document, Skill or test case?

An editorial working template based on Day 2 methods.

26 / REFERENCES

Sources and replay links

Based on the Day 2 broadcast. These public links let readers revisit the source material.

Day 2 broadcast replay

The broadcast was later titled Grok Bot builds a Game Studio LIVE.

Grok Bot Galaxy official event and schedule

Day 2 schedule; planned times are distinct from actual broadcast times.

Sales Engineering | Amrita Venkatraman

Customer cases, technical work and competitor research.

Sales | Krista Letz and Mark Wright

Research, meetings, CRM and customer follow-up.

SDRs | Simon Lackowski

Prospecting, priorities and delegation.

Customer Support | David Gan

Knowledge, support workflows and improvement.

The Morning Newspaper | Karen X. Cheng

Official template page corresponding to the guest demonstration.

Time and scope

All times use SGT (UTC+8). Day 2 refers to 16 September on the US West Coast; local viewing crosses from the evening of 16 September into the morning of 17 September. Timeline entries approximate actual segments. Case figures, personal experiences and live prototypes are not official guarantees.

These are session notes, not a verbatim transcript. Gaps in evidence are not filled in as completed outcomes. HR applications and the working template are marked as editorial extensions.